

1. Data sources and gating

- Required DB tables:
 - proposal, proposer, insured_member, documents, optionally previous_insurance_details, and rule_engine_trail.
- Gating conditions:
 - documents.validated = TRUE
 - rule_engine_trail.finreview_required = TRUE for the same proposal_number
- Outcome: Only proposals meeting both conditions proceed to scoring.

2. SQL extraction (data_extraction.py)

- Uses CTEs:
 - validated_documents: proposal→proposer links with validated documents.
 - finreview_required_set: proposals marked finreview_required = TRUE.
 - member_sum_assured: SUM of insured_member.sum_insured per proposer.
 - previous_insurance_summary: SUM of other insurance per proposer (0 if table missing).
- Final SELECT returns only the fields required to score:
 - proposal_number, proposer_id, stated_age, dob, occupation, annual_income, premium, sum_assured, other_insurance_sum_assured
- Defensive behavior:
 - If optional tables don't exist (e.g., previous insurance), aggregates default to 0.
 - Schema/table names are env-overridable; code quotes identifiers and introspects columns to avoid name mismatches.

3. Preprocessing and validation (finance_score_engine.py)

- Preprocessing:
 - Safe numeric casting for annual_income, premium, sum_assured, other_insurance_sum_assured (stripping commas, negative to None).
 - Normalizes occupation to lowercased trimmed strings.
- Validation (FR2):
 - Flags missing critical fields per proposal in validation_issues (e.g., missing_annual_income).
 - These flags appear in the scored JSON for auditability and can be enforced by your consumers.

4. Ratios (derived features)

- $\text{sar_income_ratio} = \text{sum_assured} / \text{annual_income}$
- $\text{tsar_income_ratio} = (\text{sum_assured} + \text{other_insurance_sum_assured}) / \text{annual_income}$

- $\text{premium_income_ratio} = \text{premium} / \text{annual_income}$
- All ratios guard against division by zero or missing values (result becomes None if not computable).

5. YAML rule configuration (finance_score_rules.yaml)

- Weights (must sum to 1.0):
 - $\text{sar_income_ratio}: 0.5$
 - $\text{tsar_income_ratio}: 0.25$
 - $\text{premium_income_ratio}: 0.25$
- Bands per component:
 - Each ratio has ordered, non-overlapping ranges that map to 1–5 scores.
 - Example (sar): $<2 \rightarrow 1, 2-3 \rightarrow 2, 3-4 \rightarrow 3, 4-5 \rightarrow 4, \geq 5 \rightarrow 5$.
- Decisions:
 - risk_categories : maps final score 1→Safe ... 5→Reject.
 - $\text{underwriting_flags}$: conditional rules based on final score and component scores (e.g., score 2 with $\text{premium_score} \leq 4 \rightarrow \text{Pass}$; 5 → Decline).

6. Scoring logic

- Component scores:
 - Each ratio is looked up against its YAML bands to get a component score 1–5.
- Weighted total:
 - $\text{weighted_score} = \text{sar_score} \times \text{w_sar} + \text{tsar_score} \times \text{w_tsar} + \text{premium_score} \times \text{w_prem}$.
 - $\text{final_finance_score} = \text{round}(\text{weighted_score})$.
- Decisions:
 - risk_category from score→label mapping (YAML).
 - underwriting_flag from conditional rules (YAML).
- Explainability:
 - score_factors : top three contributors computed via component score × weight, sorted desc.
 - This shows why a score ended up where it did (e.g., SAR contributed most).

7. Outputs (per proposal)

- Inputs JSON (traceability):
 - Path: `finance_scores/YYYYMMDD/inputs/finance_input.json`
 - Fields: `proposal_number`, `proposer_id`, `stated_age`, `dob`, `occupation`, `annual_income`, `premium`, `sum_assured`, `other_insurance_sum_assured`
- Score JSON (decision package):
 - Path: `finance_scores/YYYYMMDD/finance_score.json`

- Fields: proposal_number, proposer_id, ratios, component scores, final_finance_score, risk_category, underwriting_flag, score_factors, validation_issues

8. Guarantees and safeguards

- Only proposals satisfying document validation and finreview_required = TRUE are processed.
- Safe handling of missing or malformed numeric fields.
- Optional tables handled with zero-value fallbacks to keep scoring operational.
- Full auditability via per-proposal inputs snapshot and scored output, plus validation_issues and score_factors.

Summary

- Extracts only proposals that are both validated and flagged for finance review.
- Computes three ratios from extracted fields (SAR, TSAR, Premium vs Income).
- Scores each ratio using YAML-defined bands, weights them, and rounds to a final score.
- Derives risk category and underwriting flag from YAML decisions.
- Writes two JSONs per proposal: inputs (for traceability) and scored outputs (for decisions).